

# *Compressing Data via Dimensionality Reduction*

**Aprendizaje Automático**

Ingeniería de Robótica Software

Universidad Rey Juan Carlos

# Content

- Introduction
- Feature extraction
- Principal Component Analysis (PCA)
- Linear Discriminant Analysis (LDA)
- t-distributed Stochastic Neighbor Embedding (t-SNE)

# Introduction

- **Different methods** exist for **reducing the dimensionality of datasets**.
- **Feature selection techniques** are one approach to achieve this.
- An alternative to feature selection for dimensionality reduction is **feature extraction**, that involves transforming the dataset into a new feature subspace with lower dimensionality.

# Introduction

- **Dimensionality reduction** or data compression is **important** in some cases in **machine learning**.
- **Why?**
  - By transforming high-dimensional data into a **lower-dimensional space**, the **complexity of the dataset is reduced**, which can result in **faster training times** and **less use of computing resources**.
  - By **discarding irrelevant or noisy data** helps to create more **robust models**.
  - By reducing the number of dimensions, it enables a **visualization of the data** (2 or 3 dimensions), which facilitates the identification of patterns and clusters.

# Introduction

- The difference between feature selection and feature extraction is that:
  - In **feature selection**, we **select a subset of the original features** from the dataset. The goal is to **identify** and keep the **most relevant features** that contribute to the predictive power of the model.
  - In **feature extraction**, we **transform** or project the original dataset onto a **new feature space**. This involves **creating new features** by **combining the existing ones**, resulting in a reduced dimensionality that can **captures the essential information** in the dataset.

# Feature extraction

- **Feature extraction** can be understood as an approach to **data compression** with the goal of **maintaining most of the relevant information**.
- In practice, **feature extraction** is used to:
  - Improve **storage space**.
  - Improve the **computational efficiency** of the learning.
  - Enable a **visualization of the data** (2D or 3D).
  - Improve the predictive performance by **discarding irrelevant or noisy**
  - Reduce the **curse of dimensionality**.

# Feature extraction


- What **problems** are caused by a high dimensionality (the **curse of dimensionality**)?
  - In **high dimensions**, data tends to become sparser. This means that **data points are farther apart** (most of the high-dimensional space is empty), **making it difficult to identify patterns** (making clustering and classification tasks challenging).
  - So, in **high dimensions it is required more data** to fill the empty space and obtain meaningful results.
  - If you **don't have more data**, **algorithms are prone to overfitting** (they fail to generalize correctly). In an attempt to capture all the variability in the dataset, the **model** can become **complex, fitting to irrelevant details that do not generalize well to new data**.

# Principal Component Analysis (PCA)

- Principal Component Analysis (**PCA**) is an **unsupervised linear transformation technique** widely used across different fields for feature extraction and dimensionality reduction.
- PCA aims to **find the directions of maximum variance in high-dimensional data** and **projects the data onto a new subspace** with equal or **fewer dimensions** than the original one.
- The **orthogonal axes** (principal components) of the **new subspace can be interpreted as the directions of maximum variance** given the constraint that the new feature axes are orthogonal to each other.




# Principal Component Analysis (PCA)



*Figure 5.1: Using PCA to find the directions of maximum variance in a dataset*

$x_1$  and  $x_2$  are the original feature axes, and PC 1 and PC 2 are the principal components.

# Principal Component Analysis (PCA)



$$\mathbf{x}\mathbf{W} = \mathbf{z}$$

**$\mathbf{W}$  is a  $d \times k$ -dimensional transformation matrix** that allows us to **map a  $d$ -dimensional vector of the features of the training example,  $\mathbf{x}$ , onto a new  $k$ -dimensional feature subspace** that has fewer dimensions than the original  $d$ -dimensional feature space.

Typically,  $k \ll d$

# Principal Component Analysis (PCA)

- As a result of **transforming the original  $d$ -dimensional dataset onto this new  $k$ -dimensional subspace**, the **first principal component** will have the **largest possible variance**.
- All **consequent principal components** will have the **largest variance given the constraint that these components are uncorrelated (orthogonal)** to the other principal components.
- **Even if the input features are correlated**, the **resulting principal components** will be **mutually orthogonal (uncorrelated)**.

# Principal Component Analysis (PCA)

- **PCA** directions are highly **sensitive to data scaling**, and **we need to standardize the features prior to PCA** if the features were measured on different scales and we want to **assign equal importance to all features**.
- **Standardization** typically involves:
  - **Centering the data: Subtracting the mean** of each feature from the dataset so that each feature has a **mean of zero**.
  - **Scaling the data: Dividing each feature by its standard deviation** so that each feature has a **standard deviation of one**.

This process is often referred to as **z-score normalization**.

By standardizing the data, you ensure that each **feature contributes equally to the calculation of the principal components**, allowing PCA to identify the directions of maximum variance without being biased by the scale of the features.

# Principal Component Analysis (PCA)

- Let's summarize the approach in a few simple steps:
  1. **Standardize** the  $d$ -dimensional dataset.
  2. Construct the **covariance matrix**.
  3. **Decompose** the **covariance matrix** into its **eigenvectors** and **eigenvalues**.
  4. **Sort the eigenvalues by decreasing order** to rank the corresponding **eigenvectors**.
  5. **Select  $k$  eigenvectors**, which correspond to the  $k$  largest eigenvalues, where  $k$  is the dimensionality of the new feature subspace ( $k \leq d$ ).
  6. **Construct a projection matrix**,  $W$ , from the “top”  $k$  eigenvectors.
  7. **Transform the  $d$ -dimensional input dataset**,  $X$ , using the projection matrix,  $W$ , to obtain the new  $k$ -dimensional feature subspace.

# Principal Component Analysis (PCA)

La eigendecomposition de una matriz  $A$  se expresa como:

$$A = V\Lambda V^{-1}$$

donde:

- $V$  es una matriz cuyas columnas son los autovectores de  $A$ .
- $\Lambda$  es una matriz diagonal cuyos elementos son los autovalores de  $A$ .
- $V^{-1}$  es la inversa de la matriz  $V$ .

- The **covariance matrix** is a special case of a **square matrix**: it's a **symmetric** matrix, which means that the matrix is equal to its transpose,  $A = A^T$ .
- When we decompose (**eigendecomposition**) such a symmetric matrix, the **eigenvalues are real** (rather than complex) **numbers**, and the **eigenvectors are orthogonal** (perpendicular) **to each other**.
- Furthermore, eigenvalues and eigenvectors come in pairs. If we decompose a covariance matrix into its eigenvectors and eigenvalues, **the eigenvectors associated with the highest eigenvalue corresponds to the direction of maximum variance in the dataset**.

# Principal Component Analysis (PCA)

```
>>> import pandas as pd
>>> df_wine = pd.read_csv(
...     'https://archive.ics.uci.edu/ml/'
...     'machine-learning-databases/wine/wine.data',
...     header=None
... )
```

Split dataset in train set  
and test set

```
>>> from sklearn.model_selection import train_test_split
>>> X, y = df_wine.iloc[:, 1:].values, df_wine.iloc[:, 0].values
>>> X_train, X_test, y_train, y_test = \
...     train_test_split(X, y, test_size=0.3,
...                       stratify=y,
...                       random_state=0)
>>> # standardize the features
>>> from sklearn.preprocessing import StandardScaler
>>> sc = StandardScaler()
>>> X_train_std = sc.fit_transform(X_train)
>>> X_test_std = sc.transform(X_test)
```

Standardize the  $d$ -  
dimensional dataset.

# Principal Component Analysis (PCA)

- **After standardization**, the **covariance matrix** is constructed.
- It is a symmetric  $d \times d$ -dimensional matrix, where  $d$  is the number of features in the dataset.
- It **stores the pairwise covariances between the different features**.
- For example, the covariance between two features can be calculated via the following equation:

$$\sigma_{jk} = \frac{1}{n-1} \sum_{i=1}^n (x_j^{(i)} - \mu_j)(x_k^{(i)} - \mu_k)$$

means of features  $j$  and  $k$

Note that, in our case, the means are zero because we standardized the dataset.



# Principal Component Analysis (PCA)

- For example, the covariance matrix of **three features** can then be written as follows:

$$\Sigma = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_2^2 & \sigma_{23} \\ \sigma_{31} & \sigma_{32} & \sigma_3^2 \end{bmatrix}$$

- The eigenvectors of the covariance matrix represent the principal components (the directions of maximum variance), whereas the corresponding eigenvalues will define their magnitude

$$A = V\Lambda V^{-1}$$

$$AV = \Lambda V$$



$$\Sigma V = \Lambda V$$

# Principal Component Analysis (PCA)

Compute the  
covariance  
matrix

```
>>> import numpy as np
>>> cov_mat = np.cov(X_train_std.T)
>>> eigen_vals, eigen_vecs = np.linalg.eig(cov_mat)
>>> print('\nEigenvalues \n', eigen_vals)
Eigenvalues
[ 4.84274532  2.41602459  1.54845825  0.96120438  0.84166161
  0.6620634   0.51828472  0.34650377  0.3131368   0.10754642
  0.21357215  0.15362835  0.1808613 ]
```

We use the `linalg.eig` function from NumPy to compute the eigenvectors and eigenvalues.

# Principal Component Analysis (PCA)

- Since we want to **reduce the dimensionality** of our dataset, we only **select the** subset of the **eigenvectors (principal components)** that contains **most of the information (variance)**.
- The **eigenvalues** define the magnitude of the **eigenvectors**, so we have to **sort the eigenvalues by decreasing magnitude**.
  - We are interested in the **top  $k$  eigenvectors** based on the values of their corresponding eigenvalues.

# Principal Component Analysis (PCA)

- Let's plot the **explained variance ratios** of the eigenvalues. The variance explained ratio of an eigenvalue is simply the **fraction of an eigenvalue and the total sum of the eigenvalues**:

$$\text{Explained variance ratio} = \frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$

```
>>> tot = sum(eigen_vals)
>>> var_exp = [(i / tot) for i in
...             sorted(eigen_vals, reverse=True)]
>>> cum_var_exp = np.cumsum(var_exp)
>>> import matplotlib.pyplot as plt
>>> plt.bar(range(1,14), var_exp, align='center',
...         label='Individual explained variance')
>>> plt.step(range(1,14), cum_var_exp, where='mid',
...          label='Cumulative explained variance')
>>> plt.ylabel('Explained variance ratio')
>>> plt.xlabel('Principal component index')
>>> plt.legend(loc='best')
>>> plt.tight_layout()
>>> plt.show()
```




Figure 5.2: The proportion of the total variance captured by the principal components

# Principal Component Analysis (PCA)

- Using the NumPy cumsum function, we can then calculate the cumulative sum of explained variances.
- The resulting plot indicates that the first principal component alone accounts for approximately 40 % of the variance in the dataset.
- Also, we can see that the first two principal components combined explain almost 60 % of the variance in the dataset.

# Principal Component Analysis (PCA)

- We have decomposed the covariance matrix into eigenpairs (eigenvectors and eigenvalues).
- **Now:**
  - Sort the eigenvalues by decreasing order.
  - Select  $k$  eigenvectors, which correspond to the  $k$  largest eigenvalues, where  $k$  is the dimensionality of the new feature subspace ( $k \leq d$ ).
  - Construct a projection matrix,  $W$ , from the  $k$  eigenvectors.
  - Transform the  $d$ -dimensional input dataset,  $X$ , using the projection matrix,  $W$ , to obtain the new  $k$ -dimensional feature subspace.

# Principal Component Analysis (PCA)

```
>>> # Make a list of (eigenvalue, eigenvector) tuples
>>> eigen_pairs = [(np.abs(eigen_vals[i]), eigen_vecs[:, i])
...                 for i in range(len(eigen_vals))]
>>> # Sort the (eigenvalue, eigenvector) tuples from high to low
>>> eigen_pairs.sort(key=lambda k: k[0], reverse=True)
```

```
>>> w = np.hstack((eigen_pairs[0][1][:, np.newaxis],
...                 eigen_pairs[1][1][:, np.newaxis]))
>>> print('Matrix W:\n', w)
```

Matrix W:


```
[[-0.13724218  0.50303478]
 [ 0.24724326  0.16487119]
 [-0.02545159  0.24456476]
 [ 0.20694508 -0.11352904]
 [-0.15436582  0.28974518]]
```

2 dimensions




# Principal Component Analysis (PCA)

1 sample


$$x' = xW$$

```
>>> X_train_std[0].dot(w)
array([ 2.38299011,  0.45458499])
```

The whole dataset


$$X' = XW$$


```
>>> X_train_pca = X_train_std.dot(w)
```



# Principal Component Analysis (PCA)

- Let's visualize the transformed training dataset in a 2-dimensional scatterplot.

```
>>> colors = ['r', 'b', 'g']
>>> markers = ['o', 's', '^']
>>> for l, c, m in zip(np.unique(y_train), colors, markers):
...     plt.scatter(X_train_pca[y_train==l, 0],
...                 X_train_pca[y_train==l, 1],
...                 c=c, label=f'Class {l}', marker=m)
>>> plt.xlabel('PC 1')
>>> plt.ylabel('PC 2')
>>> plt.legend(loc='lower left')
>>> plt.tight_layout()
>>> plt.show()
```



Although we encoded the class label information for the purpose of illustration in the preceding scatterplot, we have to keep in mind **that PCA is an unsupervised technique** that doesn't use any class label information.

# Principal Component Analysis (PCA)

- Principal component analysis in scikit-learn

```
>>> from sklearn.linear_model import LogisticRegression
>>> from sklearn.decomposition import PCA
>>> # initializing the PCA transformer and
>>> # logistic regression estimator:
>>> pca = PCA(n_components=2)
>>> lr = LogisticRegression(multi_class='ovr',
...                          random_state=1,
...                          solver='lbfgs')
>>> # dimensionality reduction:
>>> X_train_pca = pca.fit_transform(X_train_std)
>>> X_test_pca = pca.transform(X_test_std)
>>> # fitting the logistic regression model on the reduced dataset:
>>> lr.fit(X_train_pca, y_train)
>>> plot_decision_regions(X_train_pca, y_train, classifier=lr)
>>> plt.xlabel('PC 1')
>>> plt.ylabel('PC 2')
>>> plt.legend(loc='lower left')
>>> plt.tight_layout()
>>> plt.show()
```

# Principal Component Analysis (PCA)

- Principal component analysis in scikit-learn




Figure 5.4: Training examples and logistic regression decision regions after using scikit-learn's PCA for dimensionality reduction




Figure 5.5: Test datapoints with logistic regression decision regions in the PCA-based feature space

```
>>> plot_decision_regions(X_test_pca, y_test, classifier=lr)
>>> plt.xlabel('PC 1')
>>> plt.ylabel('PC 2')
>>> plt.legend(loc='lower left')
>>> plt.tight_layout()
>>> plt.show()
```

# Principal Component Analysis (PCA)

- If we are interested in the **explained variance ratios** of the different principal components, we can simply initialize the PCA class with the `n_components` parameter set to `None`, so all principal components are kept and the explained variance ratio can then be accessed via the `explained_variance_ratio_` attribute.

```
>>> pca = PCA(n_components=None)
>>> X_train_pca = pca.fit_transform(X_train_std)
>>> pca.explained_variance_ratio_
array([ 0.36951469, 0.18434927, 0.11815159, 0.07334252,
```

# Principal Component Analysis (PCA)

- The new features represent linear combinations of the original features with the principal components.

$$X = \begin{bmatrix} x_{11} & x_{12} \\ x_{21} & x_{22} \end{bmatrix} \quad W = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix}$$

$$X' = \begin{bmatrix} x_{11} & x_{12} \\ x_{21} & x_{22} \end{bmatrix} \cdot \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} = \begin{bmatrix} x_{11} \cdot w_{11} + x_{12} \cdot w_{21} & x_{11} \cdot w_{12} + x_{12} \cdot w_{22} \\ x_{21} \cdot w_{11} + x_{22} \cdot w_{21} & x_{21} \cdot w_{12} + x_{22} \cdot w_{22} \end{bmatrix}$$

Linear combination



Each  $w_{ij}$  is the **contributions of the original features to the new feature**.

# Principal Component Analysis (PCA)

- How ever, to measure the **contributions of the original features to the new feature**, we use a scaled versions of  $W$ , where each eigenvector is multiplied by the square root of its eigenvalue.
- The result is often called **loadings**.
- Why to scale?
  - The **resulting values can then be interpreted as the correlation between the original features and the new features**.

```
>>> loadings = eigen_vecs * np.sqrt(eigen_vals)
```

# Principal Component Analysis (PCA)

```
>>> fig, ax = plt.subplots()
>>> ax.bar(range(13), loadings[:, 0], align='center')
>>> ax.set_ylabel('Loadings for PC 1')
>>> ax.set_xticks(range(13))
>>> ax.set_xticklabels(df_wine.columns[1:], rotation=90)
>>> plt.ylim([-1, 1])
>>> plt.tight_layout()
>>> plt.show()
```

Plot the loadings for the first principal component, `loadings[:, 0]`, which is the first column in this matrix.




Figure 5.6: Feature correlations with the first principal component

For example, Alcohol has a negative correlation with the first new feature (approximately  $-0.3$ ), whereas Malic acid has a positive correlation (approximately  $0.54$ ).

# Principal Component Analysis (PCA)

We can obtain the loadings from a fitted scikit-learn PCA object in a similar manner, where `pca.components_` represents the eigenvectors and `pca.explained_variance_` represents the eigenvalues.

```
>>> sklearn_loadings = pca.components_.T * np.sqrt(pca.explained_variance_)
```

```
>>> fig, ax = plt.subplots()
>>> ax.bar(range(13), sklearn_loadings[:, 0], align='center')
>>> ax.set_ylabel('Loadings for PC 1')
>>> ax.set_xticks(range(13))
>>> ax.set_xticklabels(df_wine.columns[1:], rotation=90)
>>> plt.ylim([-1, 1])
>>> plt.tight_layout()
>>> plt.show()
```



# Principal Component Analysis (PCA)




Figure 5.7: Feature correlations to the first principal component using scikit-learn

# Principal Component Analysis (PCA)

- **Ejercicio**

# Linear Discriminant Analysis (LDA)

- After exploring PCA as an unsupervised feature extraction technique, we introduce **Linear Discriminant Analysis (LDA)**.
- LDA is a **linear transformation technique** that takes **class label information** into account (**supervised** algorithm).
  - It is a linear transformation technique that can be **used to reduce the number of dimensions**.
- LDA is sometimes also called **Fisher's LDA**.



# Linear Discriminant Analysis (LDA)

- LDA seeks to **find a linear projection** of the data that **maximizes the separability between different classes**.
- In other words, LDA attempts to **identify one or more directions in the feature space** (linear discriminants) that **maximize the distance between class means**.

# Linear Discriminant Analysis (LDA)

- The **maximum number of linear discriminants** that it can be obtained with LDA is determined by the number of classes and the number of features in your data set.
- Specifically, the maximum number of linear discriminants is the smaller of:
  - $C-1$ , where **C is the number of classes** in the data set.
  - $d$ , where **d is the number of original features**.
- Therefore, the maximum number of linear discriminants is:
  - $\min(C-1, d)$ .

# Linear Discriminant Analysis (LDA)



# Linear Discriminant Analysis (LDA)

- The let's briefly summarize the **main steps** that are required to perform LDA:
  1. **Standardize** the d-dimensional dataset (d is the number of features).
  2. For **each class**, compute the d-dimensional **mean vector**.
  3. Construct the **between-class scatter matrix**,  $S_B$ , and the **within-class scatter matrix**,  $S_W$ .
  4. Compute the **eigenvectors** and corresponding **eigenvalues** of the **matrix**  $S_W^{-1}S_B$ .
  5. **Sort the eigenvalues by decreasing order** to rank the corresponding eigenvectors.
  6. **Choose the k eigenvectors that correspond to the k largest eigenvalues** to construct a d×k-dimensional transformation matrix, **W**; the eigenvectors are the columns of this matrix.
  7. **Project the examples onto the new feature subspace** using the transformation matrix, W.

# Linear Discriminant Analysis (LDA)

- We construct the within-class scatter matrix and between-class scatter matrix.
- Each **mean vector**,  $m_i$ , stores the mean feature value,  $\mu_m$ , **with respect to the examples of class i**:

$$m_i = \frac{1}{n_i} \sum_{x \in D_i} x_m$$
$$m_i = \begin{bmatrix} \mu_{i,alcohol} \\ \mu_{i,malic\ acid} \\ \vdots \\ \mu_{i,proline} \end{bmatrix}^T$$



# Linear Discriminant Analysis (LDA)


```
>>> np.set_printoptions(precision=4)
>>> mean_vecs = []
>>> for label in range(1,4):
...     mean_vecs.append(np.mean(
...         X_train_std[y_train==label], axis=0))
...     print(f'MV {label}: {mean_vecs[label - 1]}\n')
```

# Linear Discriminant Analysis (LDA)

- We can now compute the **within-class scatter matrix**,  $\mathbf{S}_W$ :

$$\mathbf{S}_W = \sum_{i=1}^c \mathbf{S}_i$$

$$\mathbf{S}_i = \sum_{x \in D_i} (x - m_i)(x - m_i)^T$$



Individual scatter matrices,  $\mathbf{S}_i$ , of each individual class  $i$ .

# Linear Discriminant Analysis (LDA)

```
>>> d = 13 # number of features
>>> S_W = np.zeros((d, d))
>>> for label, mv in zip(range(1, 4), mean_vecs):
...     class_scatter = np.zeros((d, d))
...     for row in X_train_std[y_train == label]:
...         row, mv = row.reshape(d, 1), mv.reshape(d, 1)
...         class_scatter += (row - mv).dot((row - mv).T)
...     S_W += class_scatter
>>> print('Within-class scatter matrix: '
...       f'{S_W.shape[0]}x{S_W.shape[1]}')
Within-class scatter matrix: 13x13
```

# Linear Discriminant Analysis (LDA)

- The **assumption** that we are making when we are computing the scatter matrices is that the **class labels in the training dataset are uniformly distributed**.
- **However**, if we print the number of class labels, we see that **this assumption is violated**:

```
>>> print('Class label distribution:',  
...       np.bincount(y_train)[1:])  
Class label distribution: [41 50 33]
```

# Linear Discriminant Analysis (LDA)

- We want to **scale the individual scatter matrices**,  $S_i$ , before we sum them up as the scatter matrix,  $S_W$ .
- When we divide the scatter matrices by the number of class-examples,  $n_i$ , we can see that **computing the scatter matrix is in fact the same as computing the covariance matrix**,  $\Sigma_i$  (the covariance matrix is a normalized version of the scatter matrix):

```
>>> d = 13 # number of features
>>> S_W = np.zeros((d, d))
>>> for label, mv in zip(range(1, 4), mean_vecs):
...     class_scatter = np.cov(X_train_std[y_train==label].T)
...     S_W += class_scatter
>>> print('Scaled within-class scatter matrix: '
...       f'{S_W.shape[0]}x{S_W.shape[1]}')
Scaled within-class scatter matrix: 13x13
```

$$\Sigma_i = \frac{1}{n_i} S_i = \frac{1}{n_i} \sum_{x \in D_i} (x - m_i)(x - m_i)^T$$

# Linear Discriminant Analysis (LDA)

- After we compute the scaled within-class scatter matrix (or covariance matrix), we can move on to the next step and compute **the between-class scatter matrix  $S_B$** :

$$S_B = \sum_{i=1}^c n_i (\mathbf{m}_i - \mathbf{m})(\mathbf{m}_i - \mathbf{m})^T$$

$\mathbf{m}$  is the overall mean that is computed including examples from all  $c$  classes.

# Linear Discriminant Analysis (LDA)

```
>>> mean_overall = np.mean(X_train_std, axis=0)
>>> mean_overall = mean_overall.reshape(d, 1)

>>> d = 13 # number of features
>>> S_B = np.zeros((d, d))
>>> for i, mean_vec in enumerate(mean_vecs):
...     n = X_train_std[y_train == i + 1, :].shape[0]
...     mean_vec = mean_vec.reshape(d, 1) # make column vector
...     S_B += n * (mean_vec - mean_overall).dot(
...         (mean_vec - mean_overall).T)
>>> print('Between-class scatter matrix: '
...       f'{S_B.shape[0]}x{S_B.shape[1]}')
Between-class scatter matrix: 13x13
```

# Linear Discriminant Analysis (LDA)

- We perform the **eigendecomposition** on the matrix  $\mathbf{S}_W^{-1}\mathbf{S}_B$ :

```
>>> eigen_vals, eigen_vecs =\  
...     np.linalg.eig(np.linalg.inv(S_W).dot(S_B))
```

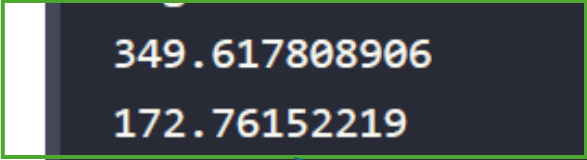
- After we compute the eigenpairs, we can sort the eigenvalues in descending order:

```
>>> eigen_pairs = [(np.abs(eigen_vals[i]), eigen_vecs[:,i])  
...                 for i in range(len(eigen_vals))]
```



# Linear Discriminant Analysis (LDA)

```
>>> eigen_pairs = sorted(eigen_pairs,  
...                       key=lambda k: k[0], reverse=True)  
>>> print('Eigenvalues in descending order:\n')  
>>> for eigen_val in eigen_pairs:  
...     print(eigen_val[0])  
Eigenvalues in descending order:  
349.617808906  
172.76152219  
3.78531345125e-14  
2.11739844822e-14
```



The number of linear discriminants is  $\min(3-1, 13) = 2$

# Linear Discriminant Analysis (LDA)

- Let's now stack the two most discriminative eigenvector columns to **create the transformation matrix,  $W$** :


```
>>> w = np.hstack((eigen_pairs[0][1][:, np.newaxis].real,  
...                eigen_pairs[1][1][:, np.newaxis].real))  
>>> print('Matrix W:\n', w)  
Matrix W:  
[[-0.1481 -0.4092]  
 [ 0.0908 -0.1577]  
 [-0.0168 -0.3537]  
 [ 0.1484  0.3223]  
 [-0.0163 -0.0817]  
 [ 0.1913  0.0842]  
 [-0.7338  0.2823]  
 [-0.075  -0.0102]]
```

# Linear Discriminant Analysis (LDA)

- Using the transformation matrix  $W$ , we can now **transform the training dataset** by multiplying the matrices:

$$X' = XW$$


```
>>> X_train_lda = X_train_std.dot(w)
>>> colors = ['r', 'b', 'g']
>>> markers = ['o', 's', '^']
>>> for l, c, m in zip(np.unique(y_train), colors, markers):
...     plt.scatter(X_train_lda[y_train==l, 0],
...                 X_train_lda[y_train==l, 1] * (-1),
...                 c=c, label= f'Class {l}', marker=m)
>>> plt.xlabel('LD 1')
>>> plt.ylabel('LD 2')
>>> plt.legend(loc='lower right')
>>> plt.tight_layout()
>>> plt.show()
```



# Linear Discriminant Analysis (LDA)

- Using scikit-learn:


```
>>> # the following import statement is one line
>>> from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
>>> lda = LDA(n_components=2)
>>> X_train_lda = lda.fit_transform(X_train_std, y_train)
>>> lr = LogisticRegression(multi_class='ovr', random_state=1,
...                          solver='lbfgs')
>>> lr = lr.fit(X_train_lda, y_train)
>>> plot_decision_regions(X_train_lda, y_train, classifier=lr)
>>> plt.xlabel('LD 1')
>>> plt.ylabel('LD 2')
>>> plt.legend(loc='lower left')
>>> plt.tight_layout()
>>> plt.show()
```



# Linear Discriminant Analysis (LDA)

- Using scikit-learn:

```
>>> X_test_lda = lda.transform(X_test_std)
>>> plot_decision_regions(X_test_lda, y_test, classifier=lr)
>>> plt.xlabel('LD 1')
>>> plt.ylabel('LD 2')
>>> plt.legend(loc='lower left')
>>> plt.tight_layout()
>>> plt.show()
```



# Linear Discriminant Analysis (LDA)

- **Ejercicio**


- En este ejercicio, deberás realizar un análisis de discriminantes lineales (LDA) utilizando el conjunto de datos proporcionado en el archivo 'dataset1.csv'.
- El análisis se realizará de dos formas diferentes: primero mediante una implementación manual (sin utilizar librerías específicas de LDA) y posteriormente usando *scikit-learn*.
- Para ambos casos, deberás calcular las discriminantes lineales que se consideren y obtener: la matriz de pesos ( $W$ ) y las proyecciones ( $X'$ ).
- Además, deberás crear visualizaciones que muestren las proyecciones ( $X'$ ) etiquetadas según su categoría.
- Finalmente, compararás los resultados obtenidos por ambos métodos.

# t-distributed Stochastic Neighbor Embedding (t-SNE)

- **PCA and LDA** are **linear transformation techniques** for feature extraction.
- However, there are **nonlinear dimensionality reduction techniques**.
  - One is t-distributed stochastic neighbor embedding (**t-SNE**), very **used to visualize high-dimensional datasets** in two or three dimensions.

# t-distributed Stochastic Neighbor Embedding (t-SNE)

- **Why** consider nonlinear dimensionality reduction?
  - Many machine learning algorithms make assumptions about the linear separability of the input data.
  - However, if **we are dealing with nonlinear problems**, which we may encounter rather frequently **in real-world applications**, linear transformation techniques for dimensionality reduction, such as PCA and LDA, may not be the best choice.





# t-distributed Stochastic Neighbor Embedding (t-SNE)


- **t-SNE** learns to **embed data points into a lower-dimensional space** such that the **pairwise distances in the original space are preserved**.
- t-SNE is a **technique intended for visualization purposes** as it **requires the whole dataset** for the projection.
- Since it projects the points directly (unlike PCA, it does **not involve a projection matrix**), **we cannot apply t-SNE to new data points**.

# t-distributed Stochastic Neighbor Embedding (t-SNE)

- We show a quick demonstration of how t-SNE can be applied to the 64-dimensional Digits dataset

```
>>> from sklearn.datasets import load_digits
>>> digits = load_digits()

>>> fig, ax = plt.subplots(1, 4)
>>> for i in range(4):
>>>     ax[i].imshow(digits.images[i], cmap='Greys')
>>> plt.show()
```



Digits are 8×8 grayscale images

```
>>> digits.data.shape
(1797, 64)
```

# t-distributed Stochastic Neighbor Embedding (t-SNE)

```
>>> y_digits = digits.target  
>>> X_digits = digits.data
```

```
>>> from sklearn.manifold import TSNE  
>>> tsne = TSNE(n_components=2, init='pca',  
...             random_state=123)  
>>> X_digits_tsne = tsne.fit_transform(X_digits)
```

# t-distributed Stochastic Neighbor Embedding (t-SNE)

```
>>> import matplotlib.patheffects as PathEffects
>>> def plot_projection(x, colors):

...     f = plt.figure(figsize=(8, 8))
...     ax = plt.subplot(aspect='equal')
...     for i in range(10):
...         plt.scatter(x[colors == i, 0],
...                     x[colors == i, 1])

...     for i in range(10):
...         xtext, ytext = np.median(x[colors == i, :], axis=0)
...         txt = ax.text(xtext, ytext, str(i), fontsize=24)
...         txt.set_path_effects([
...             PathEffects.Stroke(linewidth=5, foreground="w"),
...             PathEffects.Normal()])
```

# t-distributed Stochastic Neighbor Embedding (t-SNE)

```
>>> plot_projection(X_digits_tsne, y_digits)  
>>> plt.show()
```

